

Scalability Problems In Shared Memory

José Alves, Rui Brito

Universidade do Minho

Braga, March 2013

Index

- 1 Memory Bandwidth and Computational Intensity
- 2 Measure Task Granularity
- 3 Measure Programs Without Synchronisation
- 4 Measure Loads Per Task
- 5 Conclusion

Memory Bandwidth

$$\begin{aligned} \text{Mem_bandwidth} &= ((\text{PAPI_L2_TCM}) \times 64 \times 1 \times 10^{-9}) / \text{exec_time} \\ \text{measured_mem_bandwidth} &= (8427602 \times 64 \times 1 \times 10^{-9}) / 0.88 = \\ &0.6129 \text{ GB/s} \end{aligned}$$
$$\text{System_memory_avail_bandwidth} = 14.8473 \text{ GB/s}$$

As seen in the December presentation, conv-diff has some locality problems (has it can also be seen by the high number of L1 misses (14937146)).

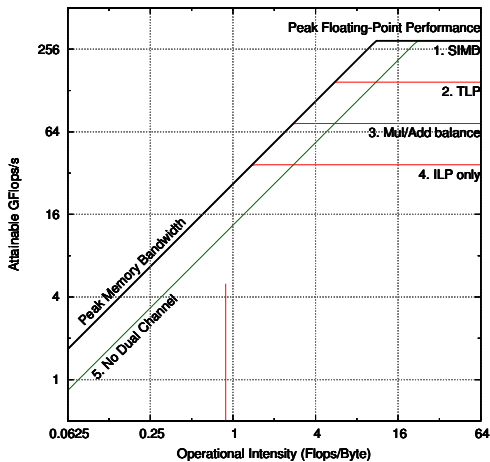


Figure : Roofline for rMBP and conv_diff

Task Granularity

In `conv_diff`:

- Only two parallel pieces of code run in parallel
- Large chunks of code
- Thread creation overhead is thus minimized

conv_diff is unsuitable for this sort of optimization:

Excessive task synchronization

- Reduction can't be used because values are updated in an array of pointers
- Synchronization must be forced on attributions

Loads Per Task

- Slight improvement using dynamic and guided scheduling
- Workload distribution also isn't the biggest problem

Conclusion

- Bad data locality hinders the entire application performance
- Implementation either AoS or SoA is expected to improve performance dramatically
- High chance that locality problem hides other problems
- Maybe what is not a problem now will prove to be so in the near future