

# Optimization of a Finite-Volume Method Application

José Alves, Rui Brito

Universidade do Minho


Braga, July 2, 2013

# Index

- 1 Introduction
- 2 Original Implementation
- 3 Tests
- 4 Optimizations
  - Naive Optimizations
  - OpenMP
  - MPI
  - Final Implementation
- 5 Results
- 6 Conclusions

# conv-diff (Recap)

- What?** Computes the heat diffusion of a fluid spreading over an area;
- How?** Uses a Finite-Volume method;
- Why?** Represents surface as a mesh, making each cell only dependent of its neighbours;



The main objective is to compute a vector  $\bar{\phi}$  such that

$$\bar{\phi} \longrightarrow G(\bar{\phi}) = \begin{pmatrix} 0 \\ 0 \\ \vdots \end{pmatrix} \text{ This is accomplished in three different stages:}$$

- ① We begin with a candidate vector  $\phi$
- ② For each edge, we compute the flux  $F_{ij}$ , with  $i$  and  $j$  being the indexes of the adjacent cells
- ③ For each cell, we compute  $\sum |e_{ij}| F_{ij} - |c_i| f_i$

$$\text{Thus: } \phi = \begin{pmatrix} \phi_1 \\ \vdots \\ \phi_I \end{pmatrix} \longrightarrow G = \begin{pmatrix} G_1 \\ \vdots \\ G_I \end{pmatrix}$$

`makeFlux` Compute the contribution from each edge;

`makeResidual` Compute the  $\phi$  vector, adding the flux for each cell from each contribution;

# Original implementation

- *Arrays-of-Pointers;*

## makeFlux

For all **edges**:

- 1 Read adjacent cell data;
- 2 Compute edge velocity;
- 3 Compute flux through edge;

## makeResidual

For all **edges**:

- 1 Subtract flux from right cell;
- 2 Add flux to left cell;

# Test Machines (for most of the project)

	compute-511@search AMD Opt 6174	compute-601@search Xeon X5650	compute-101@search Xeon E7520	MacBookPro Intel Ivy-Bridge i7
# processors	2	2	2	1
# cores per processor	12	6	1	4
hyper-threading	-	yes	yes	yes
clock frequency(GHz)	2.2	2.66	3.2	2.3
L1 capacity	128KB	32KB	16KB	64KB
L2 capacity	512KB	256KB	2MB	256KB
L3 capacity	12MB	12MB	-	6MB
RAM capacity	64GB	48GB	2GB	16GB

Table: Test machines



# Naive Optimizations

- Removed redundant loads and calculations;
- Changed some variable definitions to *const*;
- Usage of a recent compiler auto-optimizations(SLP);

# Identified Problems

- High number of memory accesses;
- Low operational intensity;
- Deep memory indirection chain;
- Bad data locality;




Figure: Total application runtime





Figure: Parallel section runtime

# Problems

- High level of communication between processes;
- High level of barrier synchronization;
- Some balancing problems;
- Computed error spikes;
- Some of FVLib's templates are hard to serialize (locality);
- Sequential portion is slow;



# Optimizations

## Array-Of-Structs

$S_1$	$e_1$
	$e_2$
	$e_3$
	...
$S_2$	$e_1$
	$e_2$
	$e_3$
	...
$S_3$	$e_1$
	$e_2$
	$e_3$
	...
...	

- Pointers  $\Rightarrow$  Indexes;

## Structs-Of-Arrays

$e_1$	$e_2$	$e_3$	...
$S_1$	$S_1$	$S_1$	
$S_2$	$S_2$	$S_2$	
$S_3$	$S_3$	$S_3$	
...	...	...	

- Pointers  $\Rightarrow$  Indexes;
- Loads only what is needed;

# OpenMP

## Implementation

- SOA;
- Similar to sequential version:
  - `parallel` for added to both core functions;


## Load Balance


- Core functions are homogeneous;
- Static scheduling:
  - Round-robin;




# Limitations

- Locality issues:
  - Softened with SOA;
  - Depends on the mesh structure;
    - gmsh does not optimize the mesh;
  - Specialized libraries:
    - parmetis;
    - Unknown complexity;







# Conclusions

- SOA proved to be the best approach;
- Parallelization did not achieve any speedups;
- Bad locality of the mesh is a limitation;
  - Execution times spike with HyperThreading;
  - Does not scale well;

# Optimization of a Finite-Volume Method Application

José Alves, Rui Brito

Universidade do Minho

Braga, July 2, 2013

– ? –